

---

# IR Modeling

Modern Information Retrieval  
by **R. Baeza-Yates and B. Ribeiro-Neto**  
(Chapter 2)

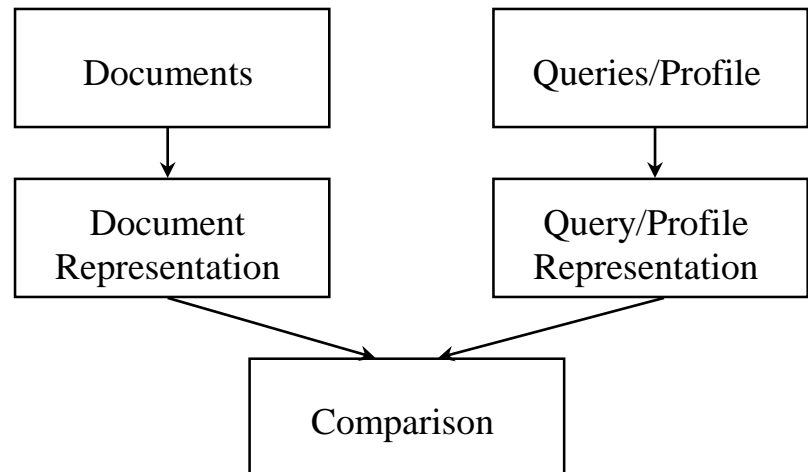
# Introduction

Remember: Judging Relevance between document and query(or profile) is important!

- » The central problem regarding IR systems is the issue of predicting which documents are relevant and which are not to a query (or a profile) by comparison.

## Taxonomy of IR Models

- » Boolean: set theoretic
- » Vector: algebraic



# Revisited: Retrievals in two ways

---

## Document Search (Ad hoc Retrieval)

- » the documents in the collection remain relatively static while new information needs are submitted to the system
- » construction of information needs as **user queries**

## Document Filtering (Routing)

- » information need remains relatively static while new documents come into the system
- » construction of the static information need as **user profile**

# Basic Concepts

## In the classic models

- » each document is described as a document vector for a set of representative keywords called index terms
- » index terms are mainly nouns and sometimes other part-of-speeches
- » distinct index terms have varying relevance weights
- » index term weights are usually assumed to be mutually independent

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

weight of index term “caesar” for “Hamlet” play

# Boolean Model (1/2)

---

Data retrieval model based on binary decision criterion

A query is a Boolean expression which can be represented as a disjunction of conjunctive vectors

- » example query:  $(\text{intelligent} \wedge \text{system}) \vee (\text{information} \wedge \text{retrieval})$

## Advantage

- » clean formalism, simplicity : explained in the next slide

## Disadvantage

- » exact matching may lead to retrieval of too few or too many documents

# Boolean Model (2/2)

## Example

doc<sub>1</sub> :


$d_1 = \{\text{"intelligent", "information", "retrieval", "learning", "agent"}\}$

doc<sub>2</sub> :

$d_2 = \{\text{"information", "management", "travel", "agent", "map"}\}$

query :

$q = (\text{"intelligent"} \wedge \text{"map"}) \vee (\text{"information"} \wedge \text{"agent"} \wedge \neg \text{"travel"})$



Therefore,  $d_1$  is relevant to  $q$  but  $d_2$  is not

# Vector Model (1/6)

Index terms are assigned **non-binary** weights

Index term weights are used to compute the degree of similarity between documents and the user query

Then, retrieved documents are sorted in decreasing order.

## Definition

*For the vector model, the weight  $w_{i,j}$  is associated with term  $k_i$  and document  $d_j$  (or query  $q$ )*

$$\vec{d}_j = (w_{1,j}, w_{2,j}, \dots, w_{t,j})$$

$$\vec{q} = (w_{1,q}, w_{2,q}, \dots, w_{t,q})$$

# Vector Model (2/6)

Degree of similarity

$$\begin{aligned} \text{sim}(d_j, q) &= \frac{\vec{d}_j \cdot \vec{q}}{|\vec{d}_j| \times |\vec{q}|} \\ &= \frac{\sum_{i=1}^t w_{i,j} \times w_{i,q}}{\sqrt{\sum_{i=1}^t w_{i,j}^2} \times \sqrt{\sum_{i=1}^t w_{i,q}^2}} = \cos \theta \end{aligned}$$




Figure 2.4 The cosine of  $\theta$  is adopted as  $\text{sim}(d_j, q)$



# Vector Model (3/6)

## Salton

- » IR can be defined as clustering a document collection into a relevant subcollection and a irrelevant one
- » Intra-cluster similarity: *tf* factor (term frequency)
- » inter-cluster dissimilarity: *idf* factor (inverse document frequency)

## Definition (1/2)

- » Term frequency is based on a simple assumption “Frequent terms are more informative than rare terms **within a document**”

$freq_{i,j}$  = number of appearances of term  $k_i$  in document  $d_j$

- » Normalized term frequency                      and                      Log term frequency

$$f_{i,j} = \frac{freq_{i,j}}{\max_l freq_{l,j}}$$

$$w_{i,j} = \begin{cases} 1 + \log_{10} freq_{i,j}, & \text{if } freq_{i,j} > 0 \\ 0, & \text{otherwise} \end{cases}$$

# Vector Model (4/6)

## Definition

- » Inverse document frequency is based on a simple assumption “Rare terms are more informative than frequent terms across documents”

$$idf_i = \log \frac{N}{n_i}$$

$N$  = total number of documents  
 $n_i$  = number of documents containing a term  $k_i$

- » document term-weighting scheme

$$w_{i,j} = freq_{i,j} \times idf_i \quad \text{often called } \textcolor{red}{tfidf} \text{ scheme}$$

- » query-term weighting scheme

$$w_{i,q} = \left(0.5 + \frac{0.5 freq_{i,q}}{\max_l freq_{l,q}}\right) \times \log \frac{N}{n_i}$$

# Vector Model (5/6)

---

## Advantages

- » its term-weighting scheme improves retrieval performance
- » its partial matching strategy allows retrieval of documents that approximate the query conditions
- » its cosine ranking formula sorts the documents according to their degree of similarity to the query

## Disadvantage

- » The assumption of mutual independence between index terms may be unrealistic in practice

# Vector Model (6/6)

Example: when using **tfidf** scheme

doc<sub>1</sub> :  $d_1 = \{\text{"intelligent"}^2, \text{"information"}, \text{"agent"}^2\}$

doc<sub>2</sub> :  $d_2 = \{\text{"information"}^2, \text{"travel"}^3, \text{"agent"}\}$

doc<sub>3</sub> :  $d_3 = \{\text{"intelligent"}, \text{"mobile"}^3, \text{"robot"}^3\}$

query :  $q = \{\text{"mobile"}, \text{"agent"}\}$

a sequence of all terms = <intelligent,information,agent,travel,mobile,robot>

$$\vec{d}_1 = (w_{1,1}, \dots, w_{6,1}) = (2 \times \log \frac{3}{2}, 1 \times \log \frac{3}{2}, 2 \times \log \frac{3}{2}, 0 \times \log \frac{3}{1}, 0 \times \log \frac{3}{1}, 0 \times \log \frac{3}{1})$$

$$\vec{d}_2 = (w_{1,2}, \dots, w_{6,2}) = (0 \times \log \frac{3}{2}, 2 \times \log \frac{3}{2}, 1 \times \log \frac{3}{2}, 3 \times \log \frac{3}{1}, 0 \times \log \frac{3}{1}, 0 \times \log \frac{3}{1})$$

$$\vec{d}_3 = (w_{1,3}, \dots, w_{6,3}) = (1 \times \log \frac{3}{2}, 0 \times \log \frac{3}{2}, 0 \times \log \frac{3}{2}, 0 \times \log \frac{3}{1}, 3 \times \log \frac{3}{1}, 3 \times \log \frac{3}{1})$$

$$\vec{q} = (w_{1,q}, \dots, w_{6,q}) = ((0.5 + \frac{0.5 \times 0}{1}) \times \log \frac{3}{2}, (0.5 + \frac{0.5 \times 0}{1}) \times \log \frac{3}{2}, (0.5 + \frac{0.5 \times 1}{1}) \times \log \frac{3}{2}, \\ (0.5 + \frac{0.5 \times 0}{1}) \times \log \frac{3}{1}, (0.5 + \frac{0.5 \times 1}{1}) \times \log \frac{3}{1}, (0.5 + \frac{0.5 \times 0}{1}) \times \log \frac{3}{1})$$

# Vector Model (6/6)

Example (Continued)

$$\vec{d}_1 = (0.35, 0.18, 0.35, 0, 0, 0)$$

$$\vec{d}_2 = (0, 0.35, 0.18, 1.43, 0, 0)$$

$$\vec{d}_3 = (0.18, 0, 0, 0, 1.43, 1.43)$$

$$\vec{q} = (0.09, 0.09, 0.18, 0.24, 0.48, 0.24)$$

$$\text{sim}(d_1, q) = \frac{\vec{d}_1 \cdot \vec{q}}{|\vec{d}_1| \times |\vec{q}|} = \dots$$

$$\text{sim}(d_2, q) = \frac{\vec{d}_2 \cdot \vec{q}}{|\vec{d}_2| \times |\vec{q}|} = \dots$$

$$\text{sim}(d_3, q) = \frac{\vec{d}_3 \cdot \vec{q}}{|\vec{d}_3| \times |\vec{q}|} = \dots$$

$$\text{sim}(d_3, q) > \text{sim}(d_1, q) > \text{sim}(d_2, q)$$

$\therefore \text{sim}(d_3, q)$  is largest, so  $d_3$  is most relevant to  $q$  !

---

For more detailed explanation of  
term frequency, document  
frequency, and cosine similarity...

# Term frequency $tf$

---

The term frequency  $tf_{t,d}$  of term  $t$  in document  $d$  is defined as the number of times that  $t$  occurs in  $d$ .

We want to use  $tf$  when computing query-document match scores. But how?

Raw term frequency is not what we want:

- » A document with 10 occurrences of the term is more relevant than a document with one occurrence of the term.
- » But not 10 times more relevant.

Relevance does not increase proportionally with term frequency.

# Log term frequency weighting

Sometimes we can use the log frequency weight of term  $t$  in  $d$

$$w_{t,d} = \begin{cases} 1 + \log_{10} \text{tf}_{t,d}, & \text{if } \text{tf}_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

The score is 0 if none of the query terms is present in the document.

$0 \rightarrow 0$ ,  $1 \rightarrow 1$ ,  $2 \rightarrow 1.3$ ,  $10 \rightarrow 2$ ,  $1000 \rightarrow 4$ , etc.

Simply, we can calculate the score for a document-query pair by just summing over terms  $t$  in both  $q$  and  $d$ :

$$\text{Simple similarity score} = \sum_{t \in q \cap d} (1 + \log \text{tf}_{t,d})$$



# Document frequency

---

Rare terms are more informative than frequent terms

Consider a term in the query that is rare in the document collection (e.g., *arachnocentric*)

A document containing this term is very likely to be relevant to the query *arachnocentric*

→ We want a high weight for rare terms like *arachnocentric*.

# Document frequency, continued

---

Consider a query term that is frequent in the document collection (e.g., *high*, *increase*, *line*)

A document containing such a term is more likely to be relevant than a document that doesn't, but it's not a sure indicator of relevance.

→ For frequent terms, we want positive weights for words like *high*, *increase*, and *line*, but lower weights than for rare terms like *arachnocentric*.

We will use document frequency (df) to capture this in the score.

df ( $\leq N$ ) is the number of documents that contain the term

# idf weight

$df_t$  is the document frequency of  $t$  that is defined as the number of documents that contain  $t$

- »  $df$  is a measure of the non-informativeness of  $t$

We define the  $idf$  (inverse document frequency) of  $t$  by

$$idf_t = \log_{10} N/df_t$$

- »  $idf$  is a measure of the informativeness of  $t$
- » We use  $\log N/df_t$  instead of  $N/df_t$  to “damping” the effect of  $idf$  because informativeness does not increase proportionally with inverse document frequency ( $N/df_t$ )

Will turn out the base of the log is immaterial.

# idf example, suppose $N = 1$ million

term	$df_t$	$idf_t$
calpurnia	1	6
animal	100	4
sunday	1,000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

There is one *idf* value for each term  $t$  in a Shakespeare collection.

# Collection vs. Document frequency

The collection frequency of  $t$  is the number of occurrences of  $t$  in the collection, counting multiple occurrences within a document.

Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

Which word is a better search term (and should get a higher weight)? Ans: Insurance Why?  $\text{Idf}_{\text{insurance}} > \text{idf}_{\text{try}}$

# tf-idf weighting

The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = \text{tf}_{t,d} \times \log_{10} N / \text{df}_t \quad \text{or}$$

$$w_{t,d} = (1 + \log \text{tf}_{t,d}) \times \log_{10} N / \text{df}_t$$

**Best known weighting scheme** in information retrieval

Note: the “-” in tf-idf is a hyphen, not a minus sign!

Alternative names: tf.idf, tfixidf

Increases with the number of occurrences within a document

Increases with the rarity of the term across documents in the collection

# Binary(Boolean) $\rightarrow$ count(tf) $\rightarrow$ weight matrix(tf $\times$ idf scheme)

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf weights  $\in \mathbb{R}^{|V|}$

# Documents as vectors

---

So we have a  $|V|$ -dimensional vector space

Terms are axes of the space

Documents are points or vectors in this space

Very high-dimensional: hundreds of millions of dimensions when you apply this to a web search engine

This is a very sparse vector - most entries are zero.



# Queries as vectors and Ranking documents

Key idea 1: Do the similar to document for queries: represent queries as vectors in the vector space using **query-term weighting scheme** :

$$w_{i,q} = (0.5 + \frac{0.5 tf_{i,q}}{\max_l tf_{l,q}}) \times \log \frac{N}{df_i} \quad \text{used for only query vector}$$

$$w_{t,d} = (1 + \log tf_{t,d}) \times \log_{10} N / df_t \quad \text{used for document and also query vectors}$$

Key idea 2: Rank documents according to their proximity to the query in this space

proximity = similarity between vectors

proximity  $\approx$  inverse of distance

Generally, rank documents in decreasing order of proximities

# Formalizing vector space proximity

---

First cut of proximity: Inverse of Euclidean distance between two vector points

- » (= Inverse of Euclidean distance between the end points of the two vectors)


Inverse of Euclidean distance is good for proximity ?

Inverse of Euclidean distance may be a bad idea. Why?

- » Because Euclidean distance is large for vectors of different lengths.

# Why distance is a bad idea

The Euclidean distance between  $\vec{q}$  and  $\vec{d}_2$  is largest (that is, their proximity is lowest) even though the distribution of terms in the query  $\vec{q}$  and the distribution of terms in the document  $\vec{d}_2$  are very similar.



# Use angle instead of Euclidean distance

Suppose that we take a document  $d$  and append it to itself. Call this document  $d'$  ( $= d \parallel d$  = two same documents).

Then, “Semantically”  $d$  and  $d'$  have the same content

The Euclidean distance between the two documents ( $d$  and  $d'$ ) can be quite large, which means their proximity is quite low when we use “inverse of distance” for proximity measure

We can use the angle between the two document vectors

- » The angle is 0 when having maximal similarity (proximity)
- » The angle is 90 when having minimum similarity (proximity)

**Key idea: Rank documents according to their angle with query in increasing order.**

# From angles to cosines

---

The following two notions are equivalent.

- » Rank documents in increasing order of the angle between query vector and document vector
- » Rank documents in decreasing order of  $\cos(\text{angle}(\text{query vector}, \text{document vector}))$

It is because *cosine* is a monotonically decreasing function for the interval  $[0^\circ, 90^\circ]$

# Before explaining cosine similarity, check Length normalization of Vector

A vector can be (length-) normalized by dividing each of its components by its length – for this we use the norm of vector:

$$\|\vec{x}\|_2 = \sqrt{\sum_i x_i^2}$$

Dividing a vector by its norm makes it a unit (length) vector

Effect on the two documents  $d$  and  $d'$  ( $d$  appended to itself) from earlier slide: they ( $d$  and  $d'$ ) have identical vectors after length-normalization. ( $\rightarrow$  trivial!)

# cosine(query vector, document vector)

The diagram illustrates the derivation of the cosine similarity formula. It features three main annotations in red text boxes:

- Dot product**: Points to the dot product  $\vec{q} \bullet \vec{d}$  in the first fraction of the formula.
- Unit vectors after normalization**: Points to the normalized vectors  $\frac{\vec{q}}{|\vec{q}|}$  and  $\frac{\vec{d}}{|\vec{d}|}$  in the second fraction.
- Length independent proximity measure**: A box at the bottom with dashed lines pointing to the denominator of the final formula, indicating that the measure is independent of vector lengths.

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \bullet \frac{\vec{d}}{|\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

$q_i$  is the tf-idf weight of term  $i$  in the query vector

$d_i$  is the tf-idf weight of term  $i$  in the document vector

$\cos(\vec{q}, \vec{d})$  is the cosine similarity of  $\vec{q}$  and  $\vec{d}$  ... or,  
equivalently, the cosine of the angle between  $\vec{q}$  and  $\vec{d}$ .